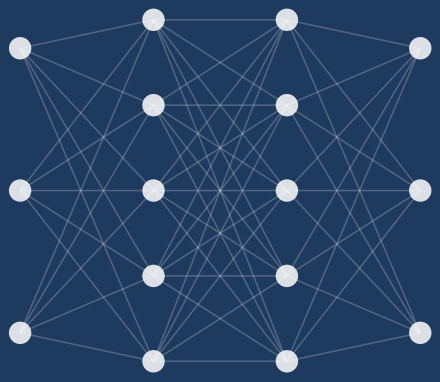
DEEP LEARNING FOUNDATIONS | PRACTICAL LABS

Practical Lab 5

Batch vs Stochastic vs Mini-batch

Dataset: Fashion-MNIST clothing images
Framework: TensorFlow / Keras



SUBMITTED BY	Arsalan
PROGRAM	Artificial Intelligence | 4th Semester
COURSE	Deep Learning Foundations
LAB	Lab 5 - Batch vs Stochastic vs Mini-batch Gradient Descent
DATE	1 November 2026

Executed code (Google Colab notebook)

https://colab.research.google.com/github/24pwai0015-max/ann-and-deep-learning/blob/main/Lab5_Batch_vs_Stochastic_vs_MiniBatch.ipynb

Click the link to open the notebook in Colab (Runtime > Run all).

1. Introduction

Gradient descent can estimate the gradient from different numbers of training examples per weight update. **Full-batch** descent uses the whole training set for each update, **stochastic** descent updates after every single example, and **mini-batch** descent uses a small group of examples (here 32). In Keras this is controlled only by the `batch_size` argument of `model.fit()`; the optimizer itself stays unchanged.

This lab keeps `SGD(learning_rate=0.02)`, the number of epochs and the initial weights fixed and changes only the batch size for three architectures: an ANN (one hidden layer), an MLP (two hidden layers) and a CNN (two convolution blocks), all trained on a small Fashion-MNIST subset.

Batch regimes compared

Technique	Batch size	Updates / epoch	Total updates (3 epochs)
Full batch	600	1	3
Stochastic	1	600	1800
Mini-batch	32	19	57

2. Experimental setup

Item	Setting
Dataset	Fashion-MNIST (28 x 28 grayscale clothing images, 10 classes)
Subset sizes	Training 600 | Validation 200 | Test 500 (fixed random subsets, seed 42)
ANN	Flatten → Dense(64, ReLU) → Dense(10, softmax)
MLP	Flatten → Dense(128, ReLU) → Dense(64, ReLU) → Dense(10, softmax)
CNN	Conv2D(16) → MaxPool → Conv2D(32) → MaxPool → Flatten → Dense(32) → Dense(10, softmax)
Fixed settings	`SGD(learning_rate=0.02)`, EPOCHS = 3, loss = sparse categorical cross-entropy, same seed
Initial weights	Saved once per model and restored before each of its three runs
Only change between runs	`batch_size` in `model.fit()`: 600, 1 or 32

3. Source code

The complete program is shown below with an explanatory comment on every statement. It is the same code that is executed in the Colab notebook **Lab5_Batch_vs_Stochastic_vs_MiniBatch.ipynb**: open in Google Colab.

```
# LAB 5: Compare full-batch, stochastic and mini-batch gradient descent.
# Run in Google Colab or locally after: pip install tensorflow numpy matplotlib
# Models: one-hidden-layer ANN, two-hidden-layer MLP and a compact CNN.
# Key control: use the SAME tf.keras.optimizers.SGD(learning_rate=0.02) everywhere.
# The gradient descent technique is selected by model.fit(..., batch_size=...).

import os  # Control the quantity of TensorFlow start-up messages.
os.environ["TF_CPP_MIN_LOG_LEVEL"] = "2"  # Reduce routine TensorFlow logging.
import csv  # Create a table that can be opened in Excel.
import math  # Calculate the number of weight updates per epoch.
import time  # Time each training setup on the current machine.
from pathlib import Path  # Make portable filenames and result folders.
import numpy as np  # Prepare reproducible subsets of image examples.
import matplotlib.pyplot as plt  # Visualize measured validation accuracy.
import tensorflow as tf  # Build and optimize all three neural networks.
```

```

SEED = 42  # Specify a reproducible dataset selection and random seed.
TRAIN_N = 600  # Keep single-example SGD feasible during a short classroom lab.
VAL_N = 200  # Reserve 200 images for validation after every epoch.
TEST_N = 500  # Use held-out test images after training each run.
EPOCHS = 3  # Use the same number of complete data passes in every run.
LEARNING_RATE = 0.02  # Hold the SGD learning rate constant across techniques.
OUT = Path("lab5_gradient_descent_results")  # Select a folder for CSV data and figures.
OUT.mkdir(parents=True, exist_ok=True)  # Create the output folder if it is missing.
tf.keras.utils.set_random_seed(SEED)  # Seed the supported Python and machine-learning RNGs.

(x_all, y_all), (x_test_all, y_test_all) = tf.keras.datasets.fashion_mnist.load_data()  # Load clothes images.
rng = np.random.default_rng(SEED)  # Set an independent reproducible NumPy random source.
chosen_ids = rng.permutation(len(x_all))[:TRAIN_N + VAL_N]  # Randomly select 800 development examples.
test_ids = rng.permutation(len(x_test_all))[:TEST_N]  # Randomly select separate test examples.
x_train = x_all[chosen_ids[:TRAIN_N]].astype("float32") / 255.0  # Scale training images to 0..1.
y_train = y_all[chosen_ids[:TRAIN_N]]  # Select the training target classes.
x_val = x_all[chosen_ids[TRAIN_N:]].astype("float32") / 255.0  # Scale validation images to 0..1.
y_val = y_all[chosen_ids[TRAIN_N:]]  # Select the validation target classes.
x_test = x_test_all[test_ids].astype("float32") / 255.0  # Scale the final held-out images.
y_test = y_test_all[test_ids]  # Select their correct class labels.
print("Train:", x_train.shape, "| Validation:", x_val.shape, "| Test:", x_test.shape)  # Confirm sizes.

def build_ann():  # Make the shallow fully connected neural network.
    return tf.keras.Sequential([  # Construct the layers in sequential order.
        tf.keras.layers.Input(shape=(28, 28)),  # Accept a 28-by-28 image.
        tf.keras.layers.Flatten(),  # Transform the image pixels into a vector.
        tf.keras.layers.Dense(64, activation="relu"),  # Add one hidden Dense layer.
        tf.keras.layers.Dense(10, activation="softmax"),  # Predict the 10 output classes.
    ])  # Complete the shallow ANN.

def build_mlp():  # Make a deeper multilayer perceptron architecture.
    return tf.keras.Sequential([  # Construct a fully connected model.
        tf.keras.layers.Input(shape=(28, 28)),  # Accept a grayscale image.
        tf.keras.layers.Flatten(),  # Flatten the input into 784 numeric features.
        tf.keras.layers.Dense(128, activation="relu"),  # Add the first hidden layer.
        tf.keras.layers.Dense(64, activation="relu"),  # Add the second hidden layer.
        tf.keras.layers.Dense(10, activation="softmax"),  # Predict the clothing classes.
    ])  # Complete the MLP.

def build_cnn():  # Build a compact convolutional classifier.
    return tf.keras.Sequential([  # Use the same framework as the other models.
        tf.keras.layers.Input(shape=(28, 28, 1)),  # Expect an explicit grayscale channel.
        tf.keras.layers.Conv2D(16, 3, activation="relu"),  # Learn basic visual features.
        tf.keras.layers.MaxPooling2D(),  # Reduce the feature-map sizes.
        tf.keras.layers.Conv2D(32, 3, activation="relu"),  # Learn higher-level patterns.
        tf.keras.layers.MaxPooling2D(),  # Downsample a second time.
        tf.keras.layers.Flatten(),  # Transform the feature maps into a vector.
        tf.keras.layers.Dense(32, activation="relu"),  # Combine the visual features.
        tf.keras.layers.Dense(10, activation="softmax"),  # Output the class probabilities.
    ])  # Complete the CNN network.

MODELS = [("ANN", build_ann), ("MLP", build_mlp), ("CNN", build_cnn)]  # List the three model designs.
TECHNIQUES = [("Full batch", TRAIN_N), ("Stochastic", 1), ("Mini-batch", 32)]  # Change batch_size only.

results = []  # Collect the results of three models times three regimes.
epoch_rows = []  # Collect per-epoch validation accuracy for every run.
for model_name, builder in MODELS:  # Evaluate each network architecture in turn.
    if model_name == "CNN":  # Convolution requires four-dimensional image tensors.
        x_tr = x_train[..., np.newaxis]  # Add one channel to the training images.
        x_va = x_val[..., np.newaxis]  # Add one channel to the validation images.
        x_te = x_test[..., np.newaxis]  # Add one channel to the testing images.
    else:  # Fully connected networks accept the original 28-by-28 images.
        x_tr, x_va, x_te = x_train, x_val, x_test  # Reuse the same images without a channel axis.
    tf.keras.utils.set_random_seed(SEED)  # Reproduce this network's original initialization.
    reference = builder()  # Construct a new, as-yet-untrained reference model.
    initial_weights = reference.get_weights()  # Save the identical starting parameters.
    del reference  # Free the no-longer-needed reference network.
    histories = {}  # Store each technique's true validation accuracy curve.
    for technique, batch_size in TECHNIQUES:  # Change only the batch size within this model.
        tf.keras.utils.set_random_seed(SEED)  # Reset randomness between independent runs.
        model = builder()  # Rebuild the very same architecture.
        model.set_weights(initial_weights)  # Reset exactly the same starting weights.
```

```
        optimizer = tf.keras.optimizers.SGD(learning_rate=LEARNING_RATE)  # Fix the optimizer AND its rate.
        model.compile(optimizer=optimizer,  # Attach the fixed SGD optimizer.
                      loss="sparse_categorical_crossentropy",  # Use the integer-label loss.
                      metrics=["accuracy"])  # Track classification accuracy.
        steps_per_epoch = math.ceil(TRAIN_N / batch_size)  # Count the expected updates per epoch.
        start = time.perf_counter()  # Begin timing the actual training run.
        history = model.fit(x_tr, y_train,  # Train on the fixed training subset.
                            validation_data=(x_va, y_val),  # Validate after every epoch.
                            epochs=EPOCHS,  # Use the same number of epochs every time.
                            batch_size=batch_size,  # The ONLY setting that changes: 600, 1 or 32.
                            shuffle=True,  # Shuffle the training data each epoch.
                            verbose=0)  # Train without verbose progress logs.
        elapsed = time.perf_counter() - start  # Record the elapsed wall-clock training time.
        test_loss, test_accuracy = model.evaluate(x_te, y_test, verbose=0)  # Evaluate on held-out examples.
        best_val = max(history.history["val_accuracy"])  # Check the strongest validation epoch.
        histories[technique] = history.history["val_accuracy"]  # Save the actual accuracy trajectory.
        results.append({  # Save the metrics together with the update counts.
            "model": model_name,  # Architecture name.
            "method": technique,  # Full batch, stochastic or mini-batch.
            "batch_size": batch_size,  # Examples per gradient update.
            "updates_per_epoch": steps_per_epoch,  # Weight updates in one epoch.
            "total_updates": EPOCHS * steps_per_epoch,  # Weight updates over all epochs.
            "best_val_accuracy": round(best_val, 4),  # Best validation accuracy over the epochs.
            "test_accuracy": round(test_accuracy, 4),  # Accuracy on the held-out test subset.
            "test_loss": round(test_loss, 4),  # Loss on the held-out test subset.
            "train_seconds": round(elapsed, 2),  # Training time on this machine.
        })  # Finish the result row.
        for epoch_no, acc in enumerate(history.history["val_accuracy"], start=1):  # Walk through epochs.
            epoch_rows.append({"model": model_name, "method": technique,  # Identify the run.
                               "epoch": epoch_no, "val_accuracy": round(acc, 4)})  # Save its accuracy.
        print(f"{model_name:3} | {technique:10} | batch={batch_size:3} | steps/epoch={steps_per_epoch:3} "
              f"| val={best_val:.3f} | test={test_accuracy:.3f}")  # Show the observations.
    plt.figure(figsize=(8, 4.5))  # Create the current network's accuracy plot.
    for method, curve in histories.items():  # Compare the three real training histories.
        plt.plot(range(1, EPOCHS + 1), curve, marker="o", label=method)  # Plot the observed accuracies.
    plt.title(f"{model_name}: validation accuracy by batch regime")  # Identify the current architecture.
    plt.xlabel("Epoch")  # Put complete dataset passes on the horizontal axis.
    plt.ylabel("Validation accuracy")  # Show validation performance on the vertical axis.
    plt.xticks(range(1, EPOCHS + 1))  # Display integer epoch tick marks.
    plt.grid(alpha=0.25)  # Make the slopes easier to compare.
    plt.legend()  # Label the full-batch, stochastic and mini-batch curves.
    plt.tight_layout()  # Avoid clipping axis annotations.
    plt.savefig(OUT / f"{model_name.lower()}_batch_methods.png", dpi=160)  # Save the current model chart.
    plt.close()  # Release this figure's memory.

with (OUT / "batch_method_results.csv").open("w", newline="", encoding="utf-8") as file:  # Open the CSV.
    writer = csv.DictWriter(file, fieldnames=list(results[0].keys()))  # Define the CSV columns.
    writer.writeheader()  # Save informative table headings.
    writer.writerows(results)  # Save all nine measured runs.
with (OUT / "epoch_validation_accuracy.csv").open("w", newline="", encoding="utf-8") as file:  # Second CSV.
    writer = csv.DictWriter(file, fieldnames=list(epoch_rows[0].keys()))  # Columns for the epoch table.
    writer.writeheader()  # Write the column names.
    writer.writerows(epoch_rows)  # Write the per-epoch validation accuracies.
print("CAUTION: Equal epochs do NOT mean equal updates when the batch size differs.")  # Key limitation.
print(f"Results and plots saved to: {OUT.resolve()}")  # Show the location of the output files.
```

4. Results (measured by running the code)

All values below were produced by executing the code above; none were typed in or estimated. The same numbers are saved by the program in `batch_method_results.csv` (columns: `model, method, batch_size, updates_per_epoch, total_updates, best_val_accuracy, test_accuracy, test_loss, train_seconds`). The best value in each group is shaded green.

Table 1. Student result table (measured)

Model	Technique	Batch size	Steps / epoch	Accuracy (test)	Best val. acc.	Seconds
ANN	Full batch	600	1	0.1780	0.1850	0.87
ANN	Stochastic	1	600	0.6020	0.6600	4.02
ANN	Mini-batch	32	19	0.6300	0.6700	1.07
MLP	Full batch	600	1	0.1380	0.1400	0.87
MLP	Stochastic	1	600	0.6160	0.6600	4.16
MLP	Mini-batch	32	19	0.6120	0.6250	0.88
CNN	Full batch	600	1	0.1340	0.1550	1.49
CNN	Stochastic	1	600	0.6220	0.6600	6.90
CNN	Mini-batch	32	19	0.3520	0.4150	1.50

Table 2. Validation accuracy after each epoch (measured)

Model	Technique	Epoch 1	Epoch 2	Epoch 3
ANN	Full batch	0.1400	0.1500	0.1850
ANN	Stochastic	0.6250	0.6450	0.6600
ANN	Mini-batch	0.5300	0.6450	0.6700
MLP	Full batch	0.0800	0.1000	0.1400
MLP	Stochastic	0.6100	0.6500	0.6600
MLP	Mini-batch	0.3350	0.5500	0.6250
CNN	Full batch	0.1350	0.1550	0.1450
CNN	Stochastic	0.5450	0.6300	0.6600
CNN	Mini-batch	0.2950	0.3250	0.4150

Accuracy comparison graphs

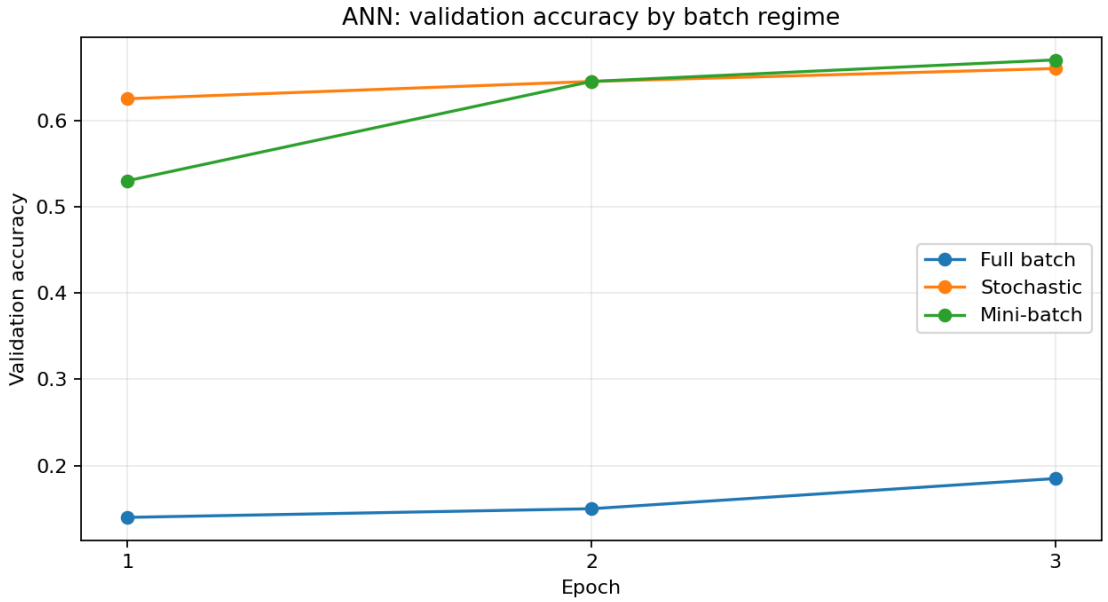


Figure 1. ANN: validation accuracy per epoch for full-batch, stochastic and mini-batch training (measured).

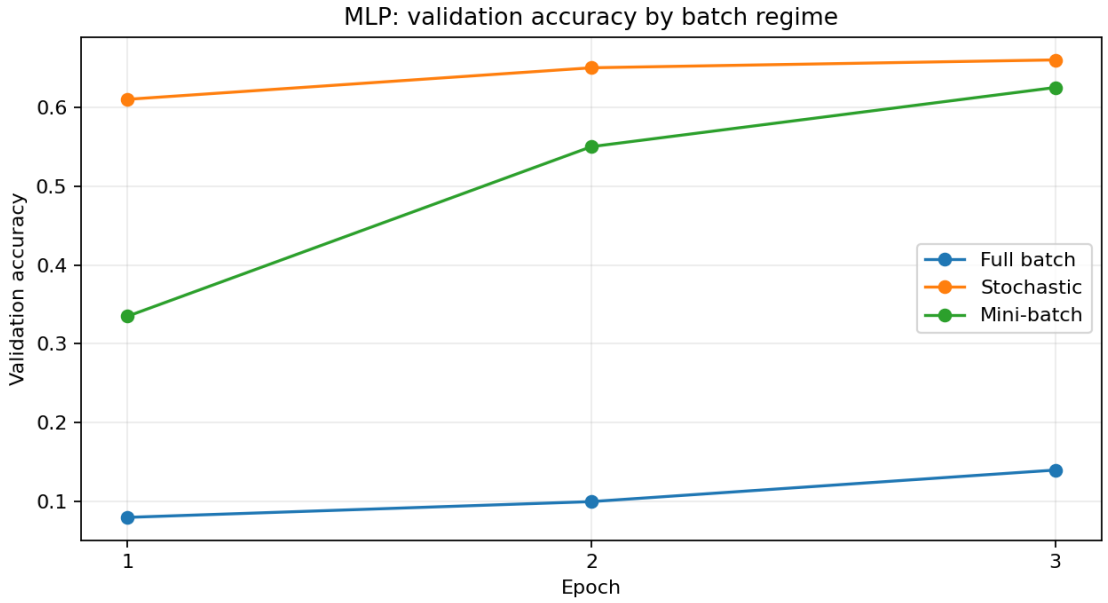


Figure 2. MLP: validation accuracy per epoch for full-batch, stochastic and mini-batch training (measured).

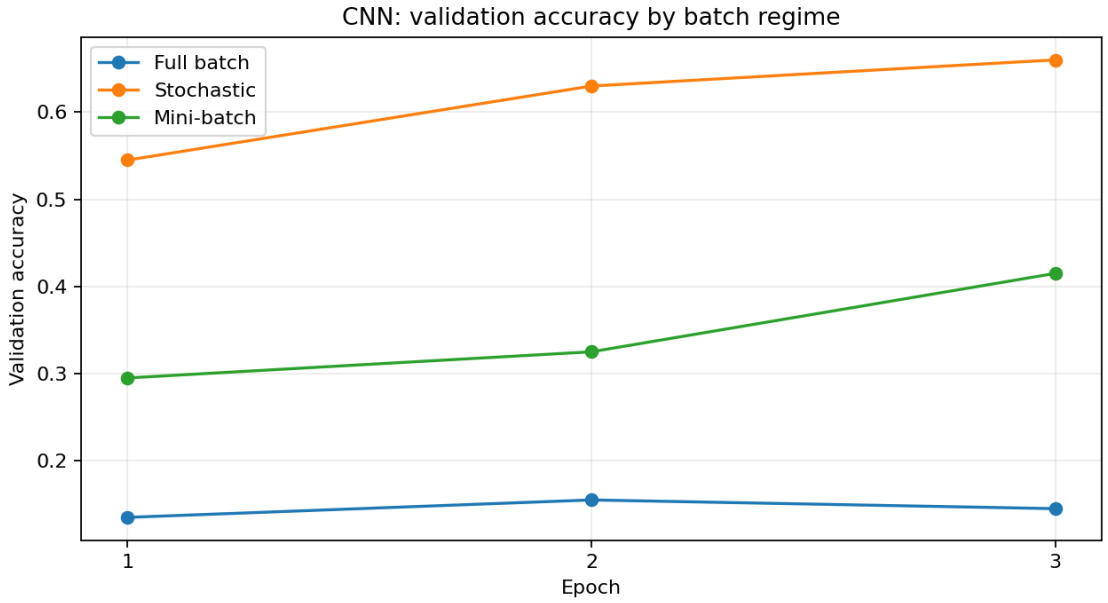


Figure 3. CNN: validation accuracy per epoch for full-batch, stochastic and mini-batch training (measured).

5. Discussion

Stochastic training (batch size 1) performs the most updates per epoch: 600, against 19 for mini-batch (600 / 32, rounded up) and only 1 for full batch, because one update is made per batch and the number of batches is the training size divided by the batch size. Over three epochs that is 1800, 57 and 3 updates. Noisy single-example gradients make the path jagged, yet stochastic training still reached 0.60 to 0.62 test accuracy for all three models. The noise is costly per update, though: for the ANN and MLP, mini-batch matched it (0.630 and 0.612) with about 30 times fewer updates and roughly a quarter of the time. Full batch stayed near chance (0.134 to 0.178) because three steps at a rate of 0.02 are far too few. Three equal-epoch runs are not a matched-update comparison: equal epochs mean equal passes over the data but very different numbers of weight updates, so accuracy differences mix the effect of batch size with the effect of update count. These are also single-seed results on 600 training images.

What changed, what stayed fixed, and limitations

- **Changed:** only `batch_size` in `model.fit()` (600, 1, 32); the architecture changes only between the three blocks (ANN, MLP, CNN).
- **Fixed:** `SGD(learning_rate=0.02)`, EPOCHS = 3, image subsets, seed, and the starting weights of each model (restored before every run).
- **Limitations:** equal epochs do not give equal updates (3, 1800 and 57), so this is not a matched-update comparison; one seed per run; a very small training set (600 images); an untuned learning rate; and training seconds measured on one CPU.

Interpretation of the results

Why do updates differ? Steps per epoch = ceil(training size / batch size): ceil(600/600) = 1, ceil(600/1) = 600 and ceil(600/32) = 19.

Which regime is fastest? Mini-batch and full batch were fastest in seconds (about 0.9 to 1.5 s). Stochastic took 4.0 to 6.9 s because 1800 tiny steps cannot use vectorized computation, so per-step overhead most likely dominates (not separately measured).

How to interpret accuracy Compare accuracy together with the number of updates. Full batch scored low because it was barely trained; the CNN's mini-batch score (0.352) is likewise low because 57 updates are not enough for it at this rate.

A stronger follow-up Match the total number of updates (or the training time), repeat with several seeds, and tune the learning rate separately for each batch size.

6. Reproducibility and environment

Python 3.12, TensorFlow 2.21.0 and Keras 3.15.1 on a CPU-only runtime (one core); random seed 42 is applied before every run. The dataset is the standard public set provided through `tf.keras.datasets`. Because hardware and library versions differ, re-running in Colab can change the *seconds* noticeably and the accuracies slightly; this report quotes only the values from the run shown here.

Files produced by the code: `batch_method_results.csv`, `epoch_validation_accuracy.csv`, and the PNG graphs shown above.

7. References

Keras optimizers reference: https://keras.io/api/optimizers/

TensorFlow Fashion-MNIST tutorial: https://www.tensorflow.org/tutorials/keras/classification

TensorFlow CNN tutorial: https://www.tensorflow.org/tutorials/images/cnn

Google Machine Learning Crash Course (hyperparameters): https://developers.google.com/machine-learning/crash-course/linear-regression/hyperparameters